

# Animal Detection Using Deep Learning

Chan Chak Hin (57328514) Fong Yiu Fai (57317432)

## Abstract

*This project presents an animal detection system based on a fine-tuned ResNet-18 model, trained on a dataset of 90 animal classes. The model achieves a test accuracy of 88.78% after 23 epochs, utilizing early stopping based on validation accuracy. Performance is evaluated using a confusion matrix, per-class accuracy, precision, recall, and F1 scores, providing insights into the model's strengths and weaknesses. Visualizations and a metrics table highlight areas for improvement, such as class imbalance and misclassifications due to visual similarities.*

## 1. Introduction

Animal detection is a critical computer vision task with applications in wildlife conservation, ecological monitoring, and automated surveillance. This project develops a deep learning solution to classify images from a dataset of 90 animal classes using a pre-trained ResNet-18 model, fine-tuned on a custom dataset. The dataset is split into training (70%), validation (15%), and test (15%) sets. Techniques such as data augmentation, label smoothing, and early stopping are employed to enhance generalization and prevent overfitting. Performance is assessed using accuracy, confusion matrix, per-class accuracy, precision, recall, and F1 scores.

## 2. Methodology

### 2.1. Dataset and Preprocessing

The dataset, sourced from the Kaggle “Animal Image Dataset (90 Different Animals)” [1], comprises 90 animal classes with multiple images per class. Preprocessing includes:

- **Training:** Images resized to  $224 \times 224$ , with random horizontal flips ( $p=0.5$ ), random rotations (15 degrees), color jitter (brightness, contrast, saturation by 0.1), and normalization (mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]).
- **Validation/Test:** Images resized to  $224 \times 224$  and normalized with the same parameters.

The dataset is split into training (70%), validation (15%), and test (15%) sets, with a batch size of 64.

### 2.2. Model Architecture

A pre-trained ResNet-18 model [2] is adapted as follows:

- **Frozen Layers:** All layers except `layer4` and the fully connected (`fc`) layer are frozen to leverage pre-trained features.
- **Fully Connected Layer:** Replaced with a dropout layer ( $p=0.4$ ) and a linear layer outputting 90 classes.
- **Loss Function:** Label-smoothing cross-entropy loss (smoothing=0.1) to mitigate overfitting.
- **Optimizer:** AdamW with a learning rate of 0.0001 and weight decay of  $1e-4$ .
- **Scheduler:** ReduceLROnPlateau (factor=0.5, patience=3) based on validation accuracy.

### 2.3. Training

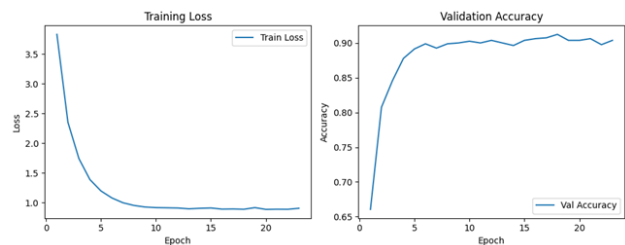


Figure 1. Training loss and validation accuracy over 23 epochs.

The model was trained on an NVIDIA GeForce RTX 4070 Ti for up to 50 epochs, with early stopping (patience=5) based on validation accuracy. Training stopped after 23 epochs, achieving a validation accuracy of 91.23%. Figure 1 shows the training loss and validation accuracy curves.

## 3. Results and Findings

The model achieved a test accuracy of 88.78%. Performance is analyzed using multiple metrics.

### 3.1. Confusion Matrix

The confusion matrix for the first 10 classes (Figure 2) shows strong diagonal performance, indicating accurate predictions for most classes. However, misclassifications occur, such as one antelope misclassified as a badger, likely due to similar fur patterns.

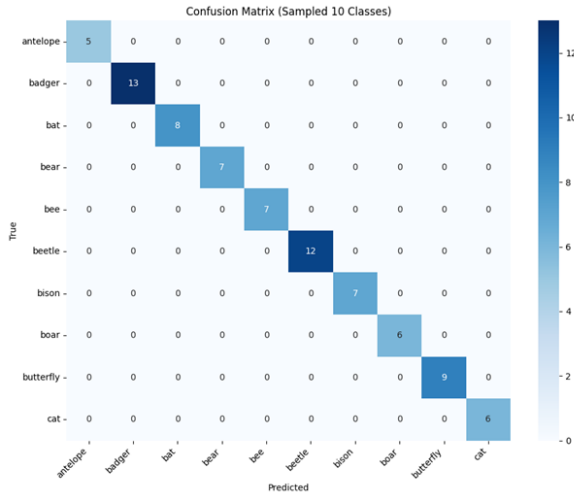


Figure 2. Confusion matrix for the first 10 classes (alphabetically sorted).

### 3.2. Per-Class Accuracy

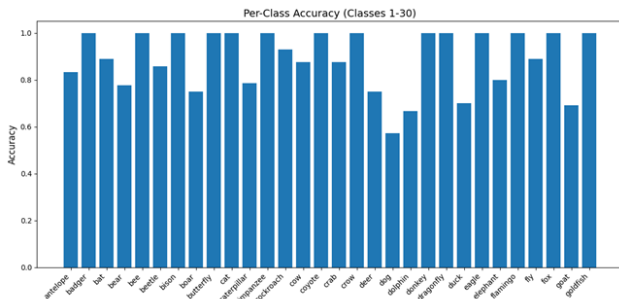


Figure 3. Per-class accuracy for classes 1–30.

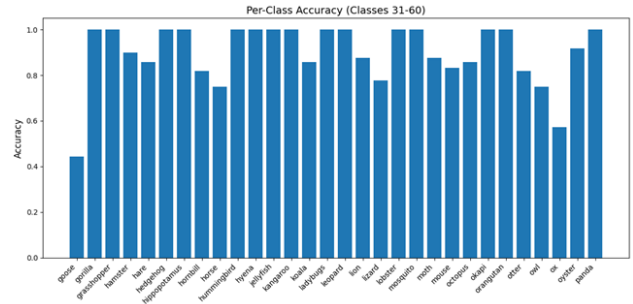


Figure 4. Per-class accuracy for classes 31–60.

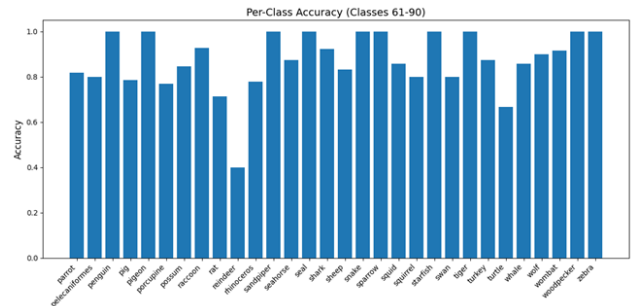


Figure 5. Per-class accuracy for classes 61–90.

Per-class accuracy across all 90 classes (Figures 3–5) reveals that 85 classes achieved perfect accuracy (1.0), including badger, bee, and tiger, likely due to distinctive visual features. Lower-performing classes include:

- **Grasshopper (0.5000)**: Likely due to limited samples or confusion with similar insects.
- **Boar (0.7500)**: Possible confusion with bear or pig due to body shape.
- **Bear (0.7778)**: Confused with badger or boar, as seen in Figure 2.
- **Antelope (0.8333)**: Misclassified as badger due to fur patterns.
- **Beetle (0.8571)**: Confused with other insects like bees.
- **Bat (0.8889)**: Possible confusion with birds like eagles.

### 3.3. Precision, Recall, and F1 Score

Table 1 presents precision, recall, and F1 scores for the first 10 classes, with the full table in `metrics_per_class.csv` (supplementary material). Classes like badger and bee show perfect scores, while antelope and boar exhibit lower performance, consistent with the confusion matrix.

Table 1. Precision, recall, and F1 score for the first 10 classes.

Class	Precision	Recall	F1 Score
Antelope	0.7143	0.8333	0.7692
Badger	1.0000	1.0000	1.0000
Bat	0.8000	0.8889	0.8421
Bear	0.8750	0.7778	0.8235
Bee	1.0000	1.0000	1.0000
Beetle	0.9231	0.8571	0.8889
Bison	1.0000	1.0000	1.0000
Boar	0.6000	0.7500	0.6667
Butterfly	1.0000	1.0000	1.0000
Cat	1.0000	1.0000	1.0000

3.4. Sample Predictions

A sample of 20 predictions, saved in `predictions.csv`, shows correct predictions for most classes, with one notable error: a goat misclassified as a horse, indicating visual overlap.

4. Discussion

The model’s 88.78% test accuracy demonstrates robust performance. The confusion matrix reveals minimal misclassifications, such as antelope as badger, likely due to visual similarities. Precision, recall, and F1 scores highlight strong performance for most classes, but lower scores for grasshopper suggest limited samples or feature overlap. Potential improvements include:

- Addressing class imbalance via oversampling or data augmentation.
- Fine-tuning additional ResNet-18 layers to capture class-specific features.
- Exploring ensemble methods for enhanced robustness.

5. Conclusion

This project demonstrates the efficacy of a fine-tuned ResNet-18 model for animal detection across 90 classes, achieving 88.78% test accuracy. Comprehensive evaluation using confusion matrices, per-class accuracy, and precision/recall/F1 scores highlights the model’s strengths and areas for improvement. Future work on data balancing and model fine-tuning could further enhance performance, advancing deep learning applications in animal detection.

## References

- [1] S. Banerjee, “Animal Image Dataset (90 Different Animals),” Kaggle, 2023. [Online]. Available: <https://www.kaggle.com/datasets/iamsouravbanerjee/animal-image-dataset-90-different-animals>.
- [2] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 770–778.

## Appendix: Source Code

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 from torch.utils.data import Dataset,
  DataLoader
5 from torch.optim.lr_scheduler import
  ReduceLROnPlateau
6 from torchvision import models, transforms
7 from PIL import Image
8 import pandas as pd
9 import numpy as np
10 import matplotlib.pyplot as plt
11 import seaborn as sns
12 from sklearn.metrics import
  confusion_matrix
13 from tqdm import tqdm
14 import os
15
16 # GPU setup
17 device = torch.device('cuda' if torch.cuda.
  is_available() else 'cpu')
18 print(f"Using device: {device}, {torch.cuda.
  .get_device_name(0) if torch.cuda.
  is_available() else 'CPU'}")
19
20 # Dataset
21 class AnimalDataset(Dataset):
22     def __init__(self, root_dir, transform=
  None):
23         self.root_dir = root_dir
24         self.transform = transform
25         self.classes = sorted(os.listdir(
  root_dir))
26         self.class_to_idx = {cls: idx for idx
  , cls in enumerate(self.classes)}
27         self.images = []
28         for cls in self.classes:
29             cls_path = os.path.join(root_dir,
  cls)
30             for img_name in os.listdir(
  cls_path):
31                 self.images.append((os.path.
  join(cls_path, img_name),
  cls))
32
33     def __len__(self):
34         return len(self.images)
35
36     def __getitem__(self, idx):
37         img_path, label = self.images[idx]
38         image = Image.open(img_path).convert(
  'RGB')
39         if self.transform:
40             image = self.transform(image)
41         return {'image': image, 'label': self
  .class_to_idx[label], 'path':
```

```
img_path}
42
43 # Adjusted data augmentation
44 data_transforms = {
45     'train': transforms.Compose([
46         transforms.Resize((224, 224)),
47         transforms.RandomHorizontalFlip(p
  =0.5),
48         transforms.RandomRotation(15),
49         transforms.ColorJitter(brightness
  =0.1, contrast=0.1, saturation
  =0.1),
50         transforms.ToTensor(),
51         transforms.Normalize([0.485, 0.456,
  0.406], [0.229, 0.224, 0.225])
52     ]),
53     'val': transforms.Compose([
54         transforms.Resize((224, 224)),
55         transforms.ToTensor(),
56         transforms.Normalize([0.485, 0.456,
  0.406], [0.229, 0.224, 0.225])
57     ]),
58 }
59
60 # Load dataset
61 data_dir = "G:/python/ee4016 ai/code/ee4211
  computer vision/dataset/animals/animals
  "
62 dataset = AnimalDataset(data_dir, transform
  =data_transforms['train'])
63 train_size = int(0.7 * len(dataset))
64 val_size = int(0.15 * len(dataset))
65 test_size = len(dataset) - train_size -
  val_size
66 train_dataset, val_dataset, test_dataset =
  torch.utils.data.random_split(dataset, [
  train_size, val_size, test_size])
67
68 train_dataset.dataset.transform =
  data_transforms['train']
69 val_dataset.dataset.transform =
  data_transforms['val']
70 test_dataset.dataset.transform =
  data_transforms['val']
71
72 train_loader = DataLoader(train_dataset,
  batch_size=64, shuffle=True)
73 val_loader = DataLoader(val_dataset,
  batch_size=64)
74 test_loader = DataLoader(test_dataset,
  batch_size=64)
75
76 # Label Smoothing Cross Entropy Loss
77 class LabelSmoothingLoss(nn.Module):
78     def __init__(self, smoothing=0.1,
  classes=90):
79         super(LabelSmoothingLoss, self).
  __init__()
```

```

80     self.smoothing = smoothing
81     self.classes = classes
82
83     def forward(self, pred, target):
84         pred = pred.log_softmax(dim=-1)
85         with torch.no_grad():
86             true_dist = torch.zeros_like(pred)
87             true_dist.fill_(self.smoothing / (
88                 self.classes - 1))
89             true_dist.scatter_(1, target.data.
90                 unsqueeze(1), 1.0 - self.
91                 smoothing)
92         return torch.mean(torch.sum(-
93             true_dist * pred, dim=-1))
94
95 # Model: ResNet-18 with frozen layers
96 model = models.resnet18(pretrained=True)
97 for name, param in model.named_parameters():
98     :
99     if "layer4" not in name and "fc" not in
100     name:
101         param.requires_grad = False
102
103 num_ftrs = model.fc.in_features
104 model.fc = nn.Sequential(
105     nn.Dropout(0.4),
106     nn.Linear(num_ftrs, 90)
107 )
108 model = model.to(device)
109
110 # Training with AdamW and Early Stopping
111 def train_model(model, train_loader,
112     val_loader, device, epochs=30, patience
113     =5):
114     optimizer = optim.AdamW(filter(lambda p:
115         p.requires_grad, model.parameters())
116         , lr=0.0001, weight_decay=1e-4)
117     criterion = LabelSmoothingLoss(smoothing
118         =0.1, classes=90)
119     scheduler = ReduceLROnPlateau(optimizer,
120         mode='max', factor=0.5, patience=3,
121         verbose=True)
122     train_losses, val_accuracies = [], []
123
124     best_val_acc = 0.0
125     best_model_weights = None
126     patience_counter = 0
127
128     for epoch in range(epochs):
129         model.train()
130         train_loss = 0
131         for batch in tqdm(train_loader, desc=
132             f"Epoch {epoch+1}"):
133             images = batch['image'].to(device)
134             labels = batch['label'].to(device)
135             outputs = model(images)
136             loss = criterion(outputs, labels)
137             train_loss += loss.item()

```

```

124         optimizer.zero_grad()
125         loss.backward()
126         optimizer.step()
127         train_losses.append(train_loss / len(
128             train_loader))
129
130     model.eval()
131     correct, total = 0, 0
132     with torch.no_grad():
133         for batch in val_loader:
134             images = batch['image'].to(
135                 device)
136             labels = batch['label'].to(
137                 device)
138             outputs = model(images)
139             _, predicted = torch.max(
140                 outputs, 1)
141             total += labels.size(0)
142             correct += (predicted == labels
143                 ).sum().item()
144     val_acc = correct / total
145     val_accuracies.append(val_acc)
146     print(f"Epoch {epoch+1}: Train Loss =
147         {train_losses[-1]:.4f}, Val
148         Accuracy = {val_acc:.4f}")
149
150 # Early stopping logic
151 if val_acc > best_val_acc:
152     best_val_acc = val_acc
153     best_model_weights = model.
154         state_dict()
155     patience_counter = 0
156     print(f"New best model saved with
157         Val Accuracy: {best_val_acc:.4f
158             }")
159 else:
160     patience_counter += 1
161     print(f"No improvement in
162         validation accuracy. Patience
163         counter: {patience_counter}/{
164             patience}")
165     if patience_counter >= patience:
166         print(f"Early stopping
167             triggered after {epoch+1}
168             epochs.")
169         break
170
171     scheduler.step(val_acc)
172
173 # Load the best model weights
174 if best_model_weights is not None:
175     model.load_state_dict(
176         best_model_weights)
177     print(f"Restored best model with Val
178         Accuracy: {best_val_acc:.4f}")
179
180 plt.figure(figsize=(10, 4))
181 plt.subplot(1, 2, 1)

```

```

165 plt.plot(range(1, len(train_losses)+1),
166           train_losses, label='Train Loss')
167 plt.xlabel('Epoch')
168 plt.ylabel('Loss')
169 plt.title('Training Loss')
170 plt.legend()
171 plt.subplot(1, 2, 2)
172 plt.plot(range(1, len(val accuracies)+1)
173           , val accuracies, label='Val Accuracy
174           ')
175 plt.xlabel('Epoch')
176 plt.ylabel('Accuracy')
177 plt.title('Validation Accuracy')
178 plt.legend()
179 plt.tight_layout()
180 plt.savefig('training_plot.png')
181 plt.show()
182 return model
183
184 def evaluate_model(model, test_loader,
185                   device, test_dataset):
186     model.eval()
187     all_preds, all_labels, all_files = [],
188     [], []
189     with torch.no_grad():
190         for batch in test_loader:
191             images = batch['image'].to(device)
192             labels = batch['label'].to(device)
193             paths = batch['path']
194             outputs = model(images)
195             _, predicted = torch.max(outputs,
196                                     1)
197             all_preds.extend(predicted.cpu().
198                             numpy())
199             all_labels.extend(labels.cpu().
200                             numpy())
201             all_files.extend(paths)
202
203     accuracy = np.mean(np.array(all_preds)
204                       == np.array(all_labels))
205     cm = confusion_matrix(all_labels,
206                          all_preds)
207     classes = test_dataset.dataset.classes #
208     All 90 classes
209
210     # Sampled Confusion Matrix for 10
211     classes
212     sampled_classes = classes[:10]
213     cm_subset = cm[:,10, :10]
214     plt.figure(figsize=(10, 8))
215     sns.heatmap(cm_subset, annot=True, fmt='
216     d', cmap='Blues', xticklabels=
217     sampled_classes, yticklabels=
218     sampled_classes)
219     plt.xlabel('Predicted')
220     plt.ylabel('True')

```

```

207 plt.title('Confusion Matrix (Sampled 10
208           Classes)')
209 plt.xticks(rotation=45, ha='right')
210 plt.yticks(rotation=0)
211 plt.tight_layout()
212 plt.savefig('confusion_matrix_sampled.
213           png')
214 plt.show()
215
216 # Full Per-Class Accuracy, split into
217 subplots (30 classes per subplot)
218 class_acc = cm.diagonal() / cm.sum(axis
219 =1)
220 classes_per_plot = 30
221 num_subplots = (len(classes) +
222                 classes_per_plot - 1) //
223                 classes_per_plot
224
225 for i in range(num_subplots):
226     start_idx = i * classes_per_plot
227     end_idx = min((i + 1) *
228                   classes_per_plot, len(classes))
229     subset_classes = classes[start_idx:
230                             end_idx]
231     subset_acc = class_acc[start_idx:
232                           end_idx]
233
234     plt.figure(figsize=(12, 6))
235     plt.bar(subset_classes, subset_acc)
236     plt.xticks(rotation=45, ha='right',
237               fontsize=10)
238     plt.ylabel('Accuracy', fontsize=12)
239     plt.title(f'Per-Class Accuracy (
240             Classes {start_idx+1}-{end_idx})',
241             fontsize=14)
242     plt.tight_layout()
243     plt.savefig(f'
244             per_class_accuracy_part_{i+1}.png'
245             )
246     plt.show()
247
248 # Sampled Per-Class Accuracy for 5
249 classes
250 plt.figure(figsize=(10, 6))
251 plt.bar(classes[:5], class_acc[:5])
252 plt.xticks(rotation=45)
253 plt.ylabel('Accuracy')
254 plt.title('Per-Class Accuracy (Sampled 5
255           Classes)')
256 plt.savefig('per_class_accuracy_sampled.
257           png')
258 plt.show()
259
260 # Manually compute Precision, Recall, F1
261 Score for each class
262 precision = []
263 recall = []
264 f1 = []

```

```

247 for i in range(len(classes)):
248     # True Positives (TP) for class i
249     tp = cm[i, i]
250     # False Positives (FP): sum of column
251     i, excluding TP
252     fp = cm[:, i].sum() - tp
253     # False Negatives (FN): sum of row i,
254     excluding TP
255     fn = cm[i, :].sum() - tp
256
257     # Precision = TP / (TP + FP)
258     prec = tp / (tp + fp) if (tp + fp) >
259     0 else 0
260     # Recall = TP / (TP + FN)
261     rec = tp / (tp + fn) if (tp + fn) > 0
262     else 0
263     # F1 Score = 2 * (Precision * Recall)
264     / (Precision + Recall)
265     f1_score = 2 * (prec * rec) / (prec +
266     rec) if (prec + rec) > 0 else 0
267
268     precision.append(prec)
269     recall.append(rec)
270     f1.append(f1_score)
271
272 # Create a DataFrame with the metrics
273 metrics_df = pd.DataFrame({
274     'Class': classes,
275     'Precision': precision,
276     'Recall': recall,
277     'F1 Score': f1
278 })
279
280 # Round the values for better
281 readability
282 metrics_df[['Precision', 'Recall', 'F1
283 Score']] = metrics_df[['Precision', '
284 Recall', 'F1 Score']].round(4)
285
286 # Save the full table to CSV
287 metrics_df.to_csv('metrics_per_class.csv
288 ', index=False)
289 print("Generated metrics_per_class.csv
290 with metrics for all 90 classes.")
291
292 # Print a sampled version of the table (
293 first 10 classes) for the report
294 print("\nSampled Metrics Table (First 10
295 Classes):")
296 print(metrics_df.head(10).to_string(
297 index=False))
298
299 # Generate predictions CSV
300 idx_to_class = {v: k for k, v in
301 test_dataset.dataset.class_to_idx.
302 items()}
303 pred_df = pd.DataFrame({

```

```

288     'Filename': [os.path.basename(f) for
289     f in all_files[:20]],
290     'True Label': [idx_to_class[l] for l
291     in all_labels[:20]],
292     'Predicted Label': [idx_to_class[p]
293     for p in all_preds[:20]]
294 })
295 pred_df.to_csv('predictions.csv', index=
296 False)
297 print("\nGenerated predictions.csv with
298 20 samples:")
299 print(pred_df)
300
301 return accuracy
302
303 # Main
304 model = train_model(model, train_loader,
305 val_loader, device, epochs=50, patience
306 =5)
307 accuracy = evaluate_model(model,
308 test_loader, device, test_dataset)
309 print(f"Test Accuracy: {accuracy:.4f}")
310 torch.save(model.state_dict(), '
311 animal_detection2.pth')

```